

Disclaimer

The "data" - For clarification, not in the scope of exam

- For this example Marelli environment has been used
- There the output of the reading operation is a custom object with several attributes and methods
- Read rdd are stored inside the ".data" attribute

Tests done with 5VM (1 master + 4 workers) for a total of 70Gb RAM and 20 Core

```
from Mco.Marc.ProcessEngine.sparkProcEng import DatabricksFullProcessEngine
from datetime import datetime, timedelta
from Mco.Marc.Base import STRINGS
from Mco.Marc.Base.abstractTracer import LOG_LEVELS
```

```
mco = DatabricksFullProcessEngine("Unimi","Join tests")
```

```
2020-11-30T08:50:40.396320Z - INFO,Unimi,Join tests,Mco.Marc.Base.abstractTracer,DefaultTracer.__init__,Log started
2020-11-30T08:50:40.473348Z - INFO,Unimi,Join tests,Mco.Marc.Base.abstractTracer,DefaultTracer.__init__,File created
2020-11-30T08:50:40.608467Z - INFO,Unimi,Join tests,Mco.Marc.ProcessEngine.sparkProcEng,DatabricksFullProcessEngine.__init__, is ready!
2020-11-30T08:50:40.670725Z - WARNING,Unimi,Join tests,Mco.Marc.ProcessEngine.sparkProcEng,DatabricksFullProcessEngine.__init__,Using Databricks Spark version with full access to primitives!
```

```

masterLookup = {}
nameLookup = {"POSNR":"Delivery_Item","MANDT":"Client","VBELN":"Delivery_Id",
              "MATNR":"Material_Number",
              "WERKS":"Receiving_Plant","LFIMG":"Quantity",
              "MEINS":"Unit_Measure","ORT01":"Customer Address",
              "VBTVV":"Document_Category","VTWEG":"Distribution_Channel",
              "SPART": "Business_Unit", "KDMAT": "Material_Id", "BWART":
              "Movement_Type", "MTART": "Mat_Type",
              "ARKTX": "Sales_Order_Text", "VKORG": "Sales_Organization",
              "KUNNR": "Customer_Id", "WADAT_IST": "Delivery_Date"}

```

```

nameLookupRdd = sc.parallelize([(key,nameLookup.get(key)) for key in
nameLookup])
nameLookupRdd.take(5)

```

```

Out[5]: [('POSNR', 'Delivery_Item'),
         ('MANDT', 'Client'),
         ('VBELN', 'Delivery_Id'),
         ('MATNR', 'Material_Number'),
         ('WERKS', 'Receiving_Plant')]

```

Data Input SAP Tables (SAP is a very popular E

```

LIPSOBJ = mco.read("/data/raw_data/global/SAP/P52/LIPS", top=10)
LIKPOBJ = mco.read("/data/raw_data/global/SAP/P52/LIKP", top=10)
KNA1OBJ = mco.read("/data/raw_data/global/SAP/P52/KNA1", top=10)

```

```

2020-11-30T08:55:35.092773Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Trying to read the path data/raw_data/global/SAP/P52/LIPS
2020-11-30T08:55:35.366925Z - ERROR,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader._read_process,Exception: Trying to read with _read_process: data/processed_data/dev/DatabricksFullProcessEngine/Unimi/Join tests/process.json
2020-11-30T08:55:36.238725Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Reading first 10 files
2020-11-30T08:56:36.312768Z - INFO,Unimi,Join tests,Mco.Marc.ProcessEngine.processEngine,DatabricksFullProcessEngine.read,639980 lines read in 61.22012209892273
2020-11-30T08:56:36.440644Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Trying to read the path data/raw_data/global/SAP/P52/LIKP
2020-11-30T08:56:36.506135Z - ERROR,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader._read_process,Exception: Trying to read with _read_process: data/processed_data/dev/DatabricksFullProcessEngine/Unimi/Join tests/process.json
2020-11-30T08:56:36.831634Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Reading first 10 files
2020-11-30T08:57:34.583191Z - INFO,Unimi,Join tests,Mco.Marc.ProcessEngine.pro

```

```
cessEngine,DatabricksFullProcessEngine.read,639980 lines read in 58.14224171638489
```

```
2020-11-30T08:57:34.637981Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Trying to read the path data/raw_data/global/SAP/P52/KNA1
```

```
2020-11-30T08:57:34.705452Z - ERROR,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader._read_process,Exception: Trying to read with _read_process: data/processed_data/dev/DatabricksFullProcessEngine/Unimi/Join tests/process.json
```

```
2020-11-30T08:57:35.047156Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Reading first 10 files
```

```
2020-11-30T08:57:51.014115Z - INFO,Unimi,Join tests,Mco.Marc.ProcessEngine.processEngine,DatabricksFullProcessEngine.read,23492 lines read in 16.37534523010254
```

```
print("LIPSObj:{LIPSObj}, \nLIKPObj:{LIKPObj}, \nKNA10bj:{KNA10bj},\nnnameLookupRdd:{nameLookupRdd}".format(LIPSObj=LIPSObj.numElem,
```

```
LIKPObj=LIKPObj.numElem,
```

```
KNA10bj=KNA10bj.numElem,
```

```
nameLookupRdd=nameLookupRdd.count()))
```

```
LIPSObj:639980,
```

```
LIKPObj:639980,
```

```
KNA10bj:23492,
```

```
nameLookupRdd:18
```

```

def keepColumn(row, columnList,ts_key=False,ts_format=False):
    newRow = {}
    for column in columnList:
        newRow[column] = row.get(column)
    if not ts_key:
        ts_key='HammerGW.ts_load'
    if not ts_format:
        ts_format=STRINGS.DATETIME_FORMAT
    newRow["ts"] = datetime.strptime(row[ts_key],ts_format)
    return newRow

def createSnapshot(obj,ts_load=False):
    snapshotObj = obj.copy()
    if not ts_load:
        ts_load = obj.header["ts_key"]
    keyList = snapshotObj.header.get("row_keys")
    getSnapshot = lambda rdd: rdd.map(lambda x: (tuple([x.get(key) for key in
keyList]),(x.get(ts_load),x)))\
        .reduceByKey(lambda x,y: x if x[0]>y[0] else y).map(lambda x: x[1]
[1])
    snapshotObj.update(getSnapshot,"getSnapshot")
    return snapshotObj

columns = ["MANDT", "VKORG", "KUNNR", "WERKS", "WADAT_IST","VBELN","TS"]
keepF = lambda rdd: rdd.map(lambda x: keepColumn(x, columns,
ts_key=False,ts_format=False))
LIKPObj.restore()
LIKPObj.update(keepF,"keepF")

smallLIKPObj = createSnapshot(LIKPObj)

2020-11-30T08:58:33.403826Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[keepF] Number of lines: 639980 in 0.005741040229797364 seconds
2020-11-30T08:58:43.218594Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[getSnapshot] Number of lines: 639980 in 0.009451396942138672 s
econds

columns = ["MANDT", "MATNR", "VTWEG", "SPART","POSNR", "WERKS", "KDMAT",
"LFIMG", "BWART", "MTART", "ARKTX", "VTWEG","MEINS","VBELN","TS","VBTVV"]
keepF = lambda rdd: rdd.map(lambda x: keepColumn(x, columns,
ts_key=False,ts_format=False))
LIPSObj.restore()
LIPSObj.update(keepF,"keepF")
smallLIPSObj = createSnapshot(LIPSObj)

```

```
2020-11-30T08:58:48.842015Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[keepF] Number of lines: 639980 in 0.005509872198104859 seconds
2020-11-30T08:58:58.039786Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[getSnapshot] Number of lines: 639980 in 0.008786808252334594 s
econds
```

```
columns = ["MANDT","KUNNR",'ORT01',"HammerGW.ts_load"]
```

```
keepF = lambda rdd: rdd.map(lambda x: keepColumn(x, columns,
ts_key=False,ts_format=False))
KNA1Obj.restore()
KNA1Obj.update(keepF,"keepF")
smallKNA1Obj = createSnapshot(KNA1Obj,ts_load="HammerGW.ts_load")
```

```
2020-11-30T08:58:59.711679Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[keepF] Number of lines: 23492 in 0.0015163912773132325 seconds
2020-11-30T08:59:03.121638Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[getSnapshot] Number of lines: 23473 in 0.003033855438232422 se
conds
```

```
smallLIPSRdd = smallLIPSOBJ.data
smallLIKPrdd = smallLIKPOBJ.data
smallKNA1rdd = smallKNA1OBJ.data
```

First Test: Two Big Rdd

```
smallLIPSRdd.first()
```

```
Out[12]: {'MANDT': '050',
'MATNR': '00519614100',
'VTWEG': ' ',
'SPART': ' ',
'POSNR': '000649',
'WERKS': 'PI36',
'KDMAT': ' ',
'LFIMG': 70.0,
'BWART': '101',
'MTART': 'ROH',
'ARKTX': '520 - DISTANZIALE BARRA STABILIZZATRICE',
'MEINS': 'PZ',
'VBELN': '0180033789',
'TS': '20190509181619',
```

```
'VBTVV': ' ',
'ts': datetime.datetime(2019, 5, 21, 3, 17, 51, 458000)}
```

Prepare Join - No Action

```
keyList = ["MANDT", "VBELN"]
plainSmallLIPSRdd = smallLIPSRdd.map(lambda x: (tuple([x.get(key) for key in
keyList]),x))
plainSmallLIKPrdd = smallLIKPrdd.map(lambda x: (tuple([x.get(key) for key in
keyList]),x))
```

```
plainSmallLIKPrdd.first()
```

```
Out[16]: (('050', '0081748519'),
{'MANDT': '050',
'VKORG': 'SI35',
'KUNNR': '0000603004',
'WERKS': ' ',
'WADAT_IST': '20170825',
'VBELN': '0081748519',
'TS': '20170825053415',
'ts': datetime.datetime(2019, 6, 4, 13, 5, 36, 157000)})
```

Direct Join - Baseline

```
plainSmallLIPSRdd.join(plainSmallLIKPrdd).count()
```

```
Out[17]: 223690
```

Lets see if Spark Cache something

```
plainSmallLIPSRdd.join(plainSmallLIKPrdd).count()
```

```
Out[18]: 223690
```

```
joined = plainSmallLIPSRdd.join(plainSmallLIKPrdd).persist()
joined.count()
```

```
Out[19]: 223690
```

```
joined.count()
```

```
Out[20]: 223690
```

Repartition

```
distributedSmallLIPsRdd = plainSmallLIPsRdd.repartition(100).persist()
distributedSmallLIKPrdd = plainSmallLIKPrdd.repartition(100).persist()
distributedSmallLIPsRdd.count(), distributedSmallLIKPrdd.count()
```

```
Out[21]: (639980, 639980)
```

Repartitioned Join

```
distributedSmallLIPsRdd.join(distributedSmallLIKPrdd).count()
```

```
Out[22]: 223690
```

Repartition + Repartitioned Join

```
distributedSmallLIPsRdd = plainSmallLIPsRdd.repartition(100).persist()
distributedSmallLIKPrdd = plainSmallLIKPrdd.repartition(100).persist()
distributedSmallLIPsRdd.count(), distributedSmallLIKPrdd.count()
distributedSmallLIPsRdd.join(distributedSmallLIKPrdd).count()
```

```
Out[14]: 223690
```

Collect as Map

```
distributedSmallLIPsMap = distributedSmallLIPsRdd.collectAsMap()
```

```
for key in distributedSmallLIPsMap.keys():
    break
```

```
key,distributedSmallLIPsMap.get(key), len(distributedSmallLIPsMap) #145.564
keys
```

```
Out[32]: (('050', '0082195107'),
{'MANDT': '050',
'MATNR': '00521138980',
'VTWEG': 'OE',
'SPART': 'CS',
'POSNR': '000660',
'WERKS': 'PI35',
'KDMAT': '00521138980',
'LFIMG': 21.0,
'BWART': '631',
'MTART': 'FERT',
'ARKTX': '520 SEMIC ANT SX',
'MEINS': 'PZ',
'VBELN': '0082195107',
'TS': '20190408092745',
```

```
'VBTVV': ' ',
'ts': datetime.datetime(2019, 5, 21, 3, 22, 11, 177000)},
145564)
```

Get

```
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallLLIPSmmap.get(x[0]),x[1])).count()
```

```
Out[16]: 639980
```

Collect As Map + Get

```
distributedSmallLLIPSmmap = distributedSmallLLIPSrdd.collectAsMap()
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallLLIPSmmap.get(x[0]),x[1])).count()
```

```
Out[17]: 639980
```

Broadcast Map

```
distributedSmallLLIPSmbroadcastMap = sc.broadcast(distributedSmallLLIPSmmap)
```

Broadcasted Map

```
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallLLIPSmbroadcastMap.value.get(x[0]),x[1])).count()
```

```
Out[34]: 639980
```

Collect + Broadcast + Broadcasted Get

```
distributedSmallLLIPSmbroadcastMap =
sc.broadcast(distributedSmallLLIPSrdd.collectAsMap())
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallLLIPSmbroadcastMap.value.get(x[0]),x[1])).count()
```

```
Out[35]: 639980
```

Second Test: One Big Rdd and a small one

Prepare Join - no Action

```
keyList = ["MANDT","KUNNR"]  
plainSmallKNA1rdd = smallKNA1rdd.map(lambda x: (tuple([x.get(key) for key in  
keyList]),x)) #23K lines  
plainSmallLIKPrdd = smallLIKPrdd.map(lambda x: (tuple([x.get(key) for key in  
keyList]),x)) #149K lines
```

Baseline - Direct Join

```
plainSmallKNA1rdd.join(plainSmallLIKPrdd).count()
```

```
Out[38]: 625568
```

Lets try to reverse the join order

```
plainSmallLIKPrdd.join(plainSmallKNA1rdd).count()
```

```
Out[39]: 625568
```

Repartition

```
distributedSmallKNA1rdd = plainSmallKNA1rdd.repartition(100).persist()  
distributedSmallLIKPrdd = plainSmallLIKPrdd.repartition(100).persist()  
distributedSmallKNA1rdd.count(), distributedSmallLIKPrdd.count()
```

```
Out[40]: (23473, 639980)
```

Repartitioned Join

```
distributedSmallKNA1rdd.join(distributedSmallLIKPrdd).count()
```

```
Out[41]: 625568
```

Repartition+Join

```
distributedSmallKNA1rdd = plainSmallKNA1rdd.repartition(100).persist()  
distributedSmallLIKPrdd = plainSmallLIKPrdd.repartition(100).persist()  
distributedSmallKNA1rdd.count(), distributedSmallLIKPrdd.count()  
distributedSmallKNA1rdd.join(distributedSmallLIKPrdd).count()
```

```
Out[26]: 625568
```

Collect As Map

```
distributedSmallKNA1map = distributedSmallKNA1rdd.collectAsMap()
```

Get instead of Join

```
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallKNA1map.get(x[0]),x[1])).count()
```

```
Out[43]: 639980
```

Collect As Map + Get

```
distributedSmallKNA1map = distributedSmallKNA1rdd.collectAsMap()
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallKNA1map.get(x[0]),x[1])).count()
```

```
Out[44]: 639980
```

Broadcast Map

```
distributedSmallKNA1SbroadcastMap = sc.broadcast(distributedSmallKNA1map)
```

Broadcasted Get

```
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallKNA1SbroadcastMap.value.get(x[0]),x[1])).count()
```

```
Out[46]: 639980
```

Map + Distribute + Broadcasted Get

```
distributedSmallKNA1SbroadcastMap =
sc.broadcast(plainSmallKNA1rdd.collectAsMap())
distributedSmallLLIKPrdd.map(lambda x:
(distributedSmallKNA1SbroadcastMap.value.get(x[0]),x[1])).count()
#5.82 sec to 0.74 sec ==> 1/3 reduction 1K€ per month ==> 300€
```

```
Out[48]: 639980
```

Third test: Big Registry

Read Data

```
bigLIPSOBJ = mco.read("/data/raw_data/global/SAP/P52/LIPS",top=40) #2.559.920
lines
bigLIKPOBJ = mco.read("/data/raw_data/global/SAP/P52/LIKP",top=40) #776.015
```

```
2020-11-30T05:43:27.487010Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Trying to read the path data/raw_data/global/SAP/P52/LIPS
2020-11-30T05:43:27.690293Z - ERROR,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader._read_process,Exception: Trying to read with _read_process: data/processed_data/dev/DatabricksFullProcessEngine/Unimi/Join tests/process.json
2020-11-30T05:43:29.002994Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Reading first 40 files
2020-11-30T05:46:10.397813Z - INFO,Unimi,Join tests,Mco.Marc.ProcessEngine.processEngine,DatabricksFullProcessEngine.read,2559920 lines read in 162.9110071659088
2020-11-30T05:46:10.558833Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Trying to read the path data/raw_data/global/SAP/P52/LIKP
2020-11-30T05:46:10.625282Z - ERROR,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader._read_process,Exception: Trying to read with _read_process: data/processed_data/dev/DatabricksFullProcessEngine/Unimi/Join tests/process.json
2020-11-30T05:46:11.042704Z - INFO,Unimi,Join tests,Mco.Marc.MarcX.reader,SparkReader.read,Reading first 40 files
2020-11-30T05:47:25.378108Z - INFO,Unimi,Join tests,Mco.Marc.ProcessEngine.processEngine,DatabricksFullProcessEngine.read,776015 lines read in 74.81969904899597
```

Clean Useless Columns

```
columns = ["MANDT", "VKORG", "KUNNR", "WERKS", "WADAT_IST", "VBELN", "TS"]
keepF = lambda rdd: rdd.map(lambda x: keepColumn(x, columns,
ts_key=False, ts_format=False))
LIKPOBJ.restore()
LIKPOBJ.update(keepF, "keepF")
```

```
smallBigLIKPOBJ = createSnapshot(bigLIKPOBJ)
```

```
columns = ["MANDT", "MATNR", "VTWEG", "SPART", "POSNR", "WERKS", "KDMAT",
"LFIMG", "BWART", "MTART", "ARKTX", "VTWEG", "MEINS", "VBELN", "TS", "VBTVV"]
keepF = lambda rdd: rdd.map(lambda x: keepColumn(x, columns,
ts_key=False, ts_format=False))
LIPSOBJ.restore()
LIPSOBJ.update(keepF, "keepF")
smallBigLIPSOBJ = createSnapshot(bigLIPSOBJ)
```

```
2020-11-30T05:47:29.467191Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.jsonL,JsonL.update,[keepF] Number of lines: 639980 in 0.003922449588775635 seconds
2020-11-30T05:48:10.379028Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
```

```
L,JsonL.update,[getSnapshot] Number of lines: 772072 in 0.04051255869865417 seconds
2020-11-30T05:48:14.787247Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[keepF] Number of lines: 639980 in 0.004343063354492188 seconds
2020-11-30T05:49:47.509455Z - INFO,Unimi,Join tests,Mco.Marc.CustomTypes.json
L,JsonL.update,[getSnapshot] Number of lines: 2559920 in 0.09234606504440307 seconds
```

Prepare Join and Make Join - Baseline

```
keyList = ["MANDT", "VBELN"]
plainBigLIPSRdd = smallBigLIPSOBJ.data.map(lambda x: (tuple([x.get(key) for key
in keyList]),x))
plainBigLIKPrdd = smallBigLIKPOBJ.data.map(lambda x: (tuple([x.get(key) for key
in keyList]),x))
plainBigLIPSRdd.join(plainBigLIKPrdd).count()
```

```
Out[41]: 2405262
```

Repartition

```
distributedSmallBigLIPSRdd = plainBigLIPSRdd.repartition(100).persist()
distributedSmallBigLIKPrdd = plainBigLIKPrdd.repartition(100).persist()
distributedSmallBigLIPSRdd.count(), distributedSmallBigLIKPrdd.count()
```

```
Out[42]: (2559920, 772072)
```

Join Repartitioned

```
distributedSmallBigLIPSRdd.join(distributedSmallBigLIKPrdd).count()
```

```
Out[43]: 2405262
```

Reversed Join Repartitioned

```
distributedSmallBigLIKPrdd.join(distributedSmallBigLIPSRdd).count()
```

```
Out[44]: 2405262
```

Collect As Map

```
distributedSmallBigLIPSMAP = distributedSmallBigLIPSRdd.collectAsMap()
```

```
org.apache.spark.SparkException: Job aborted due to stage failure: Total size of serialized results of 41 tasks (4.0 GiB) is bigger than spark.driver.maxResultSize 4.0 GiB.
```

Small Registry Collect As Map

```
distributedSmallBigLIKPmap = distributedSmallBigLIKPrdd.collectAsMap()
```

```
ConnectException: Connection refused (Connection refused)
```

Get

```
plainBigLIPsRDD.map(lambda x:  
(distributedSmallBigLIKPmap.get(x[0]),x[1])).count()
```

```
NameError: name 'plainBigLIPsRDD' is not defined
```

...

```
distributedSmallBigLIKPbroadcast = sc.broadcast(distributedSmallBigLIKPmap)
```

```
NameError: name 'distributedSmallBigLIKPmap' is not defined
```

...

```
NameError: name 'plainBigLIPsRDD' is not defined
```